

introduction to Ansible

A little background

Automation

is about taking manual processes and placing technology around them to make them repeatable. Automation is the key to speed, consistency, scalability and repeatability.

Think about car factories in the 1900 vs modern robot-automated industries!

Benefits of Automation

What does automation enable:

- Scalability
- Reliability
- Repeatability
- Consistency
- Auditability
- Security

Some widely-known configuration Management Tools

- Salt
- Puppet
- Chef
- Ansible

What is Ansible?

Ansible is a tool for:

- Configuration Management
- Deploying software
- Orchestration
- Provisioning

Ansible features

- Based on Python
- Agentless (only needs Python on remote host)
- Only requires SSH
- Push based

Idempotency

Configure systems using shell script can be simple and effective, but:

- complex logic to follow
- env variables scoping rules
- portability issues between distributions or operating system
- they are not repeatable (e.g. **idempotent**)

With Ansible we solve this problem by writing the final destination state we want to reach; the tool makes only the necessary changes.

Idempotency example #1

```
$ adduser / useradd -b -u -d -G ...  
[...other stuff...]  
$ adduser / useradd -b -u -d -G ...  
ERROR: user 'adam' already exists
```

VS

```
# Ensure the user Adam exists in the system  
- name: Add the user 'Adam' with a specific uid and a primary group of 'sudo'  
  ansible.builtin.user:  
    name: adam  
    comment: Adam Engineer  
    uid: 1077  
    group: sudo  
    createhome: yes  
    home: /home/users/adamlis
```


Idempotency example #2

SAP HANA DEPLOYMENT EXTRACT ([source](#))

from

```
echo "vm.swappiness=60" >> /etc/sysctl.d/90-sap_hana.conf
echo "kernel.msgmni=32768" >> /etc/sysctl.d/90-sap_hana.conf
...
sysctl -p /etc/sysctl.d/90-sap_hana.conf
```

to

```
- name: setting kernel tunables
  sysctl: name={{ item.name }} value={{ item.value }} state=present
          sysctl_set=yes reload=yes
  with_items:
    - { name: kernel.msgmni, value: 32768 }
    - { name: vm.swappiness, value: 60 }
```

installing Ansible

we will use a [development container](#) for our workshop:

```
$ toolbox enter
```

When you have shell access to container, installation is simple:

```
# zypper install ansible
```

This need to be run only on the "main" node. Ansible by default works via ssh connection sending/pushing commands to other machines.


Ansible Hello World

```
$ ansible -m ping localhost
localhost | SUCCESS => {
    "changed": false,
    "ping": "pong"
}
```

this means that Ansible is correctly installed and working. `-m` stands for "use this module". The `ping` module does not **change** anything on the host, it simply reply back to test the communication.

Ansible Basic Terminology

- **Task** : A single action to perform
- **Play** : A collection of tasks
- **Playbook** : YAML file containing one or more plays



PLAYBOOK EXAMPLE: INSTALL & CONFIGURE APACHE WEBSERVER

```
# begin of playbook
---
- name: first play to install and start apache
  hosts: localhost
  connection: local
  become: yes
  tasks:
    - name: install apache2 (task1)
      zypper: name=apache2 state=latest
    - name: start apache2 (task2)
      systemd:
        state: started
        name: apache2
- name: second play, includes another play from a file
  ansible.builtin.import_playbook: otherplays.yaml
# end of playbook
```

More Terminology

- **Module** : Blob of Python code which is executed to perform task
- **Inventory**: File containing hosts and groups of hosts to run tasks
- **Role**: A mechanism for reusing and organizing code in Ansible in a standard hierarchy
- **Facts**: Builtin variables related to remote systems (i.e. ipaddress, hostname, cpu, ram, disk, etc.). They are filled-in by the `setup` module which is always run by default. Let's see the facts in our machine:

```
$ ansible localhost -m setup | less
```

running Ansible

There are two ways to run ansible:

1. ad hoc

Run a single task

```
ansible <pattern> [options]
```

2. Playbook

Run multiple tasks (a *playbook*) sequentially

```
ansible-playbook <pattern> [options]
```

Inventory

Ansible inventory is the list of hosts where we want to apply our configuration. The simplest inventory is a single file with a list of hosts and groups. The default location for this file is `/etc/ansible/hosts`. You can specify a different inventory file at the command line using the `-i <path>` option or in configuration using `inventory`.

A inventory can contain many groups of hosts and associate variables to the group or at the host level.

The inventory can be made dynamic, e.g. user can provide a script that outputs list of machines (there are many already made for most cloud providers, cmdb, etc.)

SIMPLE INVENTORY EXAMPLE

```
machine-debug.example.suse.de  
another_server-1.example.suse.de
```

[virtual_machines]

```
openqa-worker1.example.suse.asia  
srv01.example.suse.asia  
srv02.example.suse.asia  
srv03.example.suse.asia  
srv04.example.suse.asia  
srv05.example.suse.asia
```

[baremetal]

```
baremetal1.example.suse.de  
baremetal2.example.suse.de
```

[asia]

```
openqa-worker1.example.suse.asia  
srv0[1-5].example.suse.asia
```

[europe]

```
machine-debug.example.suse.de  
baremetal[1-2].example.suse.de  
another_server-1.example.suse.de
```

INVENTORY EXAMPLE with some GROUP VARS

[asia]

```
openqa-worker1.example.suse.asia  
srv0[1-5].example.suse.asia
```

[europe]

```
machine-debug.example.suse.de  
baremetal[1-2].example.suse.de  
another_server-1.example.suse.de
```

[asia:vars]

```
ntp_server=time-sync-server.example.suse.com  
nfs_path="another-nfs-server.suse.asia:/folder/blabla/pckgs"
```

[europe:vars]

```
ntp_server=ntp.suse.de  
nfs_path="11.22.33.44:/folder nfs-server.suse.de:/mnt/myfolder"
```

How Ansible talks to hosts ?

By default, Ansible uses native OpenSSH, because it supports *ControlPersist* (a performance feature), Kerberos, and options in `~/.ssh/config` such as Jump Host setup.

By default, Ansible connects to all remote devices with the user name you are using on the control node. If that user name does not exist on a remote device, you can [set a different user name for the connection](#).

If you just need to do some tasks as a different user, use [privilege escalation](#):

```
- name: Ensure the httpd service is running
  service:
    name: httpd
    state: started
    become: true
```

Facts 1/2

By default, whenever you run an Ansible playbook, Ansible first gathers some information ("facts") about each host in the play.

```
$ ansible-playbook playbook.yml
PLAY [group] *****
GATHERING FACTS *****
ok: [host1]
ok: [host2]
ok: [host3]
```

Facts 2/2

Facts can be extremely helpful when you're running playbooks; you can use gathered information like host IP addresses, CPU type, disk space, operating system information, and network interface information to change when certain tasks are run, or to change certain information used in configuration files.

to see all available facts on your pc:

```
$ ansible localhost -m ansible.builtin.setup
```

Local Facts

Another way of defining host-specific facts is to place a `.fact` file in a special directory on remote hosts, `/etc/ansible/facts.d/`. These files can be either JSON or INI files, or you could use *executables/scripts* that return JSON. As an example, create the file `/etc/ansible/facts.d/settings.fact` on a remote host, with the following contents:

```
[users]
admin=jane,john
normal=jim
```

Next, use Ansible's setup module to display the new facts on the remote host:

```
$ ansible hostname -m setup -a "filter=ansible_local"
```

Beyond Ad-Hoc commands: playbooks

The Ad-Hoc `ansible` command we used till now are limited to a single task, but most of the time the activity to perform is not a single command. We can group many tasks together to form a **play**, and many plays to form a **playbook** (and again many playbooks to form a **role**). Let's say we want to install and configure a web server and start with a basic shell script:

FROM A BASIC SHELL SCRIPT

```
# Install Apache.  
zypper install -y apache2  
# Copy configuration files.  
cp my_httpd.conf /etc/apache2/httpd.conf  
# Start Apache and configure it to run at boot.  
systemctl enable apache2.service  
systemctl start apache2.service
```

to *run* this: `$ sudo ./install-apache.sh`

this script has some issues; can you spot them ?

... TO A (BAD) ANSIBLE PLAYBOOK ...

```
---
- hosts: all
  tasks:
    - name: Install Apache.
      command: zypper install -y apache2
    - name: Copy configuration files.
      command: >
        cp my_httpd.conf /etc/apache2/httpd.conf
    - name: Start Apache and configure it to run at boot.
      command: systemctl enable apache2.service
    - command: systemctl start apache2.service
```

to *run* this: `$ sudo ansible-playbook install-apache.yml`

question: why this is quite bad practice ?

... TO A BETTER ONE

```
---  
- hosts: all  
  become: yes  
  tasks:  
    - name: Install Apache.  
      package:  
        name: apache2  
        state: latest  
    - name: Copy configuration files  
      copy:  
        src: my_httpd.conf  
        dest: /etc/apache2/httpd.conf  
        owner: root  
        group: root  
        mode: 0644  
    - name: Make sure Apache is started now and at boot  
      systemd:  
        name: apache2  
        enabled: true  
        state: started
```

Register variables

There are many times that you will want to run a command, then use its return code, stderr, or stdout to determine whether to run a later task. For these situations, Ansible allows you to use register to store the output of a particular command in a variable at runtime.

- name: Run a shell command and register its output as a variable
ansible.builtin.shell: /usr/bin/foo
register: foo_result
ignore_errors: true
- name: Run a shell command using output of the previous task
ansible.builtin.shell: /usr/bin/bar
when: foo_result.rc == 5

When to quote variables (a YAML gotcha)

If you start a value with `{{ foo }}`, you must quote the whole expression to create valid YAML syntax.

```
- hosts: app_servers
  vars:
    app_path: {{ base_path }}/myapp
```

You will see: `ERROR! Syntax Error while loading YAML.` If you add quotes, Ansible works correctly:

```
- hosts: app_servers
  vars:
    app_path: "{{ base_path }}/myapp"
```

if/then/else conditionals

A task can be conditionally executed with the `when:` keyword.

```
tasks:
- name: Shut down CentOS 6 and Debian 7 systems
  ansible.builtin.command: /sbin/shutdown -t now
  when: (ansible_facts['distribution'] == "CentOS" and ansible_facts['distribution_major_version'] == "6") or
        (ansible_facts['distribution'] == "Debian" and ansible_facts['distribution_major_version'] == "7")
```

[see more example on the documentation](#)

Loops / Iteration 1/2

Repeated tasks can be written as standard loops over a simple list of strings. You can define the list directly in the task or keep the values in a variable

```
- name: Add several users
  ansible.builtin.user:
    name: "{{ item }}"
    state: present
    groups: "developers"
  loop:
    - joe
    - frank
    - "{{ another_big_list_of_users }}"
```

Loops / Iteration 2/2

You can use the `until` keyword to retry a task until a certain condition is met. Here's an example:

```
- name: Retry a task until a certain condition is met
  ansible.builtin.shell: /usr/bin/foo
  register: result
  until: result.stdout.find("all systems go") != -1
  retries: 5
  delay: 10
```

for details please [see documentation](#)

Handlers

Sometimes you want a task to run only when a change is made on a machine. For example, you may want to restart a service if a task updates the configuration of that service, but not if the configuration is unchanged. Ansible uses handlers to address this use case. **Handlers are tasks that only run when notified.**

[See documentation example](#)

Beyond the basics

- delegation / local actions
- manage pauses with `wait_for` or `prompt`
- error control: `ignore_errors` / `failed_when`
- tags / filtering
- blocks
- import, include
- reboot control

Templating

When we want to refer to some variable content, introduce some logic expressions or provide a file, we can use the **Jinja2** template engine embedded in Ansible.

A template contains variables and/or expressions, which get replaced with values when a template is rendered; and tags, which control the logic of the template. The template syntax is heavily inspired by Django and Python.

- most common is `{{ }}` for Expressions (emit the template output)
- there is also `{% %}` for Statements and `{# #}` for Comments

[See template documentation](#)

Templating Example

```
$ cd ansible_examples  
$ ansible-playbook -i inventory.txt motd.yml  
$ cat /tmp/motd
```

exercise: we want to give some control to the user, who can for example change the destination file or include/exclude IPV6 addresses. How can we achieve that ?

Ansible vault : Keeping secrets secret

If you use Ansible to fully automate the provisioning and configuration of your servers, chances are you will need to use passwords or other sensitive data for some tasks, whether it's setting a default admin password, synchronizing a private key, or authenticating to a remote service.

It's better to treat passwords and sensitive data specially, and there are two primary ways to do this:

1. Use a separate secret management service, such as Vault by HashiCorp, Keywhiz by Square, or a hosted service like AWS's Key Management Service or Microsoft Azure's Key Vault.
2. Use Ansible Vault, which is built into Ansible and stores encrypted passwords and other sensitive data alongside the rest of your playbook.

Ansible Vault

How it works:

Ansible Vault works much like a real-world vault:

1. You take any YAML file you would normally have in your playbook (e.g. a variables file, host vars, group vars, role default vars, or even task includes!), and store it in the vault.
2. Ansible encrypts the vault ('closes the door'), using a key (a password you set).
3. You store the key (your vault's password) separately from the playbook in a location only you control or can access.
4. You use the key to let Ansible decrypt the encrypted vault whenever you run your playbook.

What are Ansible roles?

Roles are a way to group multiple tasks together into one container to do the automation in very effective manner with clean directory structures.

Roles are set of tasks and additional files for a certain role which allow you to break up the configurations.

It can be easily reuse the codes by anyone if the role is suitable to someone.

It can be easily modify and will reduce the syntax errors.

an example Ansible Role can be to [install a WordPress website](#). It requires a web server, php, a database, the application and some configuration

Ansible galaxy 1/2

Ansible roles are powerful and flexible; they allow you to encapsulate sets of configuration and deployable units of playbooks, variables, templates, and other files, so you can easily reuse them across different servers.

It's annoying to have to start from scratch every time, though; wouldn't it be better if people could share roles for commonly-installed applications and services? Enter [Ansible Galaxy](#).

Ansible Galaxy, or just 'Galaxy', is a repository of community-contributed Ansible content. There are thousands of roles available which can configure and deploy common applications, and they're all available through the `ansible-galaxy` command.

Ansible Galaxy 2/2

Galaxy offers the ability to add, download, and rate roles. With an account, you can contribute your own roles or rate others' roles (though you don't need an account to use roles).

```
$ ansible-galaxy role install geerlingguy.apache \
geerlingguy.mysql geerlingguy.php
```


A LAMP server in nine lines of YAML

```
# file: lamp-setup.yml
- hosts: all
  become: yes
  roles:
    - geerlingguy.mysql
    - geerlingguy.apache
    - geerlingguy.php
    - geerlingguy.php-mysql
```

```
$ ansible-playbook -i path/to/custom-inventory lamp-setup.yml
```

Anatomy of a Role

a role has only 2 mandatory subfolders:

```
role_name/  
  meta/  
  tasks/
```

If you create a directory structure like the one shown above, with a `main.yml` file in each directory, Ansible will run all the tasks defined in `tasks/main.yml` if you call the role from your playbook using the following syntax:

```
---  
- hosts: all  
  roles:  
    - role_name
```

Role scaffolding

TIP: to easily create the directory structure for a role, we can use

```
$ ansible-galaxy role init role_name
```

Running this command creates an example role in the current working directory, which you can modify to suit your needs. Using the **init** command also ensures the role is structured correctly in case you want to someday contribute the role to Ansible Galaxy.

let's look at real use case

<https://sap-linuxlab.github.io/>

TODO find some concrete and easy to follow example in the repo

Thanks!

These slides are Open Source and live in a [github repository](#), feel free to improve them ❤️

